

Big Data Analysis - Vaibhav Oberoi

1. Acquiring top 200,00 records from Stackexchange based on View Count

<Kindly refer to other attached files for codes, commands and queries>

4 queries were run one by one on data.stackexchange.com to get top 200,000 posts ordered by decreasing order of ViewCount.

2. ETL using Pig

Step 1 : We upload the 4 CSV files to the cluster

Step 2 : We make a directory in HDFS to upload the files

Hadoop fs -mkdir /Vaibhav

Step 3 : We 'Put' the files to the HDFS directory

```
vaibhavoberoi7@cluster-2265-m:~$ hadoop fs -put QueryResults*.csv /Vaibhav
```

Step 4 : Now, we open 'Pig' and load the data files using CSVExcelStorage so as to use the added functionality of YES_MULTILINE (to ignore \n) and skipping input header:

```
grunt> data2 = LOAD '/Vaibhav/QueryResults 2.csv' USING org.apache.pig.piggybank.storage.CSVExcelStorage(',', 'YES_MULTILINE','NOCHANGE','SKIP_INPUT_HEADER') AS (Id:int, PostTypeId:int, AcceptedAnswerId:int, ParentId:int, CreationDate:chararray, DeletionDate:chararray, Score:int, ViewCount:int, Body:chararray, OwnerUserId:int, OwnerDisplayName:chararray, LastEditorUserId:int, LastEditorDisplayName:chararray, LastEditDate:chararray, LastActivityDate:chararray, Title:chararray, Tags:chararray, AnswerCount:int, CommentCount:int, FavoriteCount:int, ClosedDate:chararray, CommunityOwnedDate:chararray);
2019-11-06 23:11:32,979 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - yarn.resourcemanager.system-metrics-publisher.enabled is deprecated. Instead, use yarn.system-metrics-publisher.enabled
grunt> data3 = LOAD '/Vaibhav/QueryResults 3.csv' USING org.apache.pig.piggybank.storage.CSVExcelStorage(',', 'YES_MULTILINE','NOCHANGE','SKIP_INPUT_HEADER') AS (Id:int, PostTypeId:int, AcceptedAnswerId:int, ParentId:int, CreationDate:chararray, DeletionDate:chararray, Score:int, ViewCount:int, Body:chararray, OwnerUserId:int, OwnerDisplayName:chararray, LastEditorUserId:int, LastEditorDisplayName:chararray, LastEditDate:chararray, LastActivityDate:chararray, Title:chararray, Tags:chararray, AnswerCount:int, CommentCount:int, FavoriteCount:int, ClosedDate:chararray, CommunityOwnedDate:chararray);
2019-11-06 23:12:11,451 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - yarn.resourcemanager.system-metrics-publisher.enabled is deprecated. Instead, use yarn.system-metrics-publisher.enabled
grunt> data4 = LOAD '/Vaibhav/QueryResults 4.csv' USING org.apache.pig.piggybank.storage.CSVExcelStorage(',', 'YES_MULTILINE','NOCHANGE','SKIP_INPUT_HEADER') AS (Id:int, PostTypeId:int, AcceptedAnswerId:int, ParentId:int, CreationDate:chararray, DeletionDate:chararray, Score:int, ViewCount:int, Body:chararray, OwnerUserId:int, OwnerDisplayName:chararray, LastEditorUserId:int, LastEditorDisplayName:chararray, LastEditDate:chararray, LastActivityDate:chararray, Title:chararray, Tags:chararray, AnswerCount:int, CommentCount:int, FavoriteCount:int, ClosedDate:chararray, CommunityOwnedDate:chararray);
2019-11-06 23:12:41,405 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - yarn.resourcemanager.system-metrics-publisher.enabled is deprecated. Instead, use yarn.system-metrics-publisher.enabled
grunt>
```

Step 5 : Now that we have the 4 data files, we combine them

```
grunt> combined_data = UNION data1, data2, data3, data4;
```

Step 6 : Then we shortlist columns and remove commas:

```
Details at logfile: /home/vaibhavoberoi7/pig_1573081450665.log
grunt> formatted_data = FOREACH combined_data GENERATE Id AS Id, Score AS Score, REPLACE(Body,',*', '') AS Body, OwnerUserId AS OwnerUserId, REPLACE(Title,',*', '') AS Title, REPLACE(Tags,',*', '') AS Tags;
```

Further formatting:

Step 7 : Removing \n

```
grunt> formatted_data2 = FOREACH formatted_data GENERATE Id AS Id, Score AS Score, REPLACE(Body,'\n*', '') AS Body, OwnerUserId AS OwnerUserId, REPLACE(Title,'\n*', '') AS Title, REPLACE(Tags,'\n*', '') AS Tags;
grunt>
```

Step 8 : Remove HTML tags like <p>

```
grunt> formatted_data2 = FOREACH formatted_data2 GENERATE Id AS Id, Score AS Score, REPLACE(Body,'\n*', '') AS Body, OwnerUserId AS OwnerUserId, REPLACE(Title,'\n*', '') AS Title, REPLACE(Tags,'\n*', '') AS Tags;
grunt> formatted_data3 = FOREACH formatted_data2 GENERATE Id AS Id, Score AS Score, REPLACE(Body,'<.*?>',' ') AS Body, OwnerUserId AS OwnerUserId, REPLACE(Title,'<.*?>',' ') AS Title, Tags AS Tags;
grunt>
```

Step 9: We have to Filter out null values and store the output into a new variable

valid_data = FILTER formatted_data3 BY (OwnerUserId IS NOT NULL) AND (Score IS NOT NULL);

Step 10: Finally, we Store the valid data into HDFS directory

```

grunt> valid_data = FILTER formatted_data3 BY (OwnerUserId IS NOT NULL) AND (Score IS NOT NULL);
grunt> STORE valid_data INTO '/output' USING org.apache.pig.piggybank.storage.CSVExcelStorage(',');

```

3. HIVE

Step 1 : We Create a new HDFS directory and putting our Pig Output file over there.

```

vaibhavoberoi7@cluster-2265-m:~$ ls
Output.csv pig_1573065805218.log pig_1573079423012.log pig_1573081450665.log QueryResults_1.csv QueryResults_2.csv
vaibhavoberoi7@cluster-2265-m:~$ hadoop fs -mkdir /new_dir
vaibhavoberoi7@cluster-2265-m:~$ hadoop fs -put Output.csv /new_dir
vaibhavoberoi7@cluster-2265-m:~$ hadoop fs -ls /new_dir
Found 1 items
-rw-r--r-- 2 vaibhavoberoi7 hadoop 191936463 2019-11-07 01:24 new_dir/Output.csv

```

Step 2 : Creating table data_posts to load the data from our CSV combined file.

```

hive>
> create table data_posts(Id int, Score int, Body String, OwnerUserId Int, Title String, Tags String)
> row format delimited
> FIELDS TERMINATED BY ','
> location '/new_dir';

```

Step 3 : Setting Header to True , so as to print headers in the output .set hive.cli.print.header=true

Querying with hive:

Step 4 : -- first query - 1. The top 10 posts by score

Selecting the columns id, title and score from our table, data_posts and ordering the output in a descending manner.

```

hive> SELECT id, title, score
> FROM data_posts
> ORDER BY score DESC LIMIT 10;
Query ID = vaibhavoberoi7_20191107022406_09f78b64-8b7a-41b1-9c30-2b929607014b
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1573052185904_0009)

OK
id      title      score
11227809 Why is processing a sorted array faster than processing an unsorted array? 23557
2003505 How do I delete a Git branch locally and remotely? 16071
292357 What is the difference between 'git pull' and 'git fetch'? 11369
477816 What is the correct JSON content type? 9911
231767 "What does the '"yield"' keyword do?" 9582
1642028 "What is the '"-->" operator in C++?" 8542
348170 How do I undo 'git add' before commit? 8485
6591213 How do I rename a local Git branch? 7957
79923 What and where are the stack and heap? 7816
503093 How do I redirect to another webpage? 7727
Time taken: 27.531 seconds, Fetched: 10 row(s)

```

Step 5 : -- second query - 2. The top 10 users by post score

Applying group by functionality to OwnerUserId and aggregating the column Score

```
hive>
>
> SELECT OwnerUserId, SUM(Score) AS TotalScore
> FROM data_posts
> GROUP BY OwnerUserId
> ORDER BY TotalScore DESC LIMIT 10;
Query ID = vaibhavoberoi7_20191107022721_11df0cc0-4d34-45f8-b317-9d2a577b3e40
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1573052185904_0009)

OK
owneruserid      totalscore
87234      33997
4883       24516
9951       23792
6068       22834
51816      18538
49153      17035
95592      16895
63051      16174
39677      15985
179736     14895
```

Step 6: -- third query - 3. The number of distinct users, who used the word 'hadoop' in one of their posts

Printing out the count of distinct OwnerUserId from our table data_posts in which 'Hadoop' has been mentioned.

```
hive>
>
> SELECT COUNT(DISTINCT OwnerUserId)
> FROM data_posts
> WHERE (body LIKE '%hadoop%' OR title LIKE '%hadoop%' OR tags LIKE '%hadoop%');
Query ID = vaibhavoberoi7_20191107022840_aa3a4679-754a-4c33-8e5b-a9382a7293b5
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1573052185904_0009)

OK
_c0
329
Time taken: 14.232 seconds, Fetched: 1 row(s)
```

Step 7: Creating a new directory, and saving our table in it by grouping by OwnerUserId, Body, Title.

```
hive> INSERT OVERWRITE DIRECTORY '/hiveResults'
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> SELECT OwnerUserId, Body, Title
> FROM data_posts
> GROUP BY OwnerUserId, Body, Title;
Query ID = vaibhavoberoi7_20191107023731_feb7c6ea-153d-445f-9b4f-ad91f717a9a2
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1573052185904_0010)
```

4. TF/IDF

Step 1: As we have to query TF-IDF for the top 10 users (Query 3.2), we will create a table to store its output. Creating table 'top_data' to store the output.

```
hive> create table top_data
> row format delimited
> fields terminated by ',' as
> select owneruserid, sum(score) as totalscore
> from data_posts group by owneruserid
> order by totalscore desc limit 10;
```

Step 2 : Now we extract the columns body, title and tags for these top 10 users and create a new table

Step 3 : Before proceeding to Map Reduce, we store this table's data into a csv file, as this file would be used for input to MapReduce Phase 1.

```
hive> INSERT OVERWRITE DIRECTORY '/Result_Data'
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> SELECT owneruserid, body, title
> FROM data_users
> GROUP BY owneruserid, body, title;
Query ID = vaibhavoberoi7_20191107142821_2a87397e-d0b9-40c0-b70a-5cb92aefeb5a
Total jobs = 1
```

Step 4 : We make changes to add stop words to mapper phase one file and upload Mapper&Reducer for Phases 1,2 and 3; and uploading the jar file to our cluster.

Step 5 : We run the jar file for Mapper & Reducer Phase 1 and give the Hive Stored Results file as the input

```
19/11/07 14:42:27 INFO mapreduce.Job: Running job: job_1573052185904_0015
19/11/07 14:42:35 INFO mapreduce.Job: Job job_1573052185904_0015 running in uber mode
: false
19/11/07 14:42:35 INFO mapreduce.Job: map 0% reduce 0%
19/11/07 14:42:45 INFO mapreduce.Job: map 7% reduce 0%
19/11/07 14:42:46 INFO mapreduce.Job: map 13% reduce 0%
19/11/07 14:42:48 INFO mapreduce.Job: map 33% reduce 0%
19/11/07 14:42:54 INFO mapreduce.Job: map 47% reduce 0%
19/11/07 14:42:59 INFO mapreduce.Job: map 60% reduce 0%
19/11/07 14:43:00 INFO mapreduce.Job: map 67% reduce 0%
19/11/07 14:43:02 INFO mapreduce.Job: map 73% reduce 0%
19/11/07 14:43:03 INFO mapreduce.Job: map 80% reduce 0%
19/11/07 14:43:10 INFO mapreduce.Job: map 100% reduce 0%
19/11/07 14:43:20 INFO mapreduce.Job: map 100% reduce 20%
19/11/07 14:43:21 INFO mapreduce.Job: map 100% reduce 60%
19/11/07 14:43:24 INFO mapreduce.Job: map 100% reduce 80%
19/11/07 14:43:25 INFO mapreduce.Job: map 100% reduce 100%
19/11/07 14:43:25 INFO mapreduce.Job: Job job_1573052185904_0015 completed successfully
19/11/07 14:43:25 INFO mapreduce.Job: Counters: 50
```

Step 6 : We combine the 4 result files from Step 5, upload them to HDFS and run Phase Two.

```
vaibhavoberoi7@cluster-2265-m:~$ hadoop jar /home/vaibhavoberoi7/hadoop-streaming-2.7.1.jar -file /home/vaibhavoberoi7/MapperPhaseTwo.py /home/vaibhavoberoi7/ReducerPhaseTwo.py -mapper "python MapperPhaseTwo.py" -reducer "python ReducerPhaseTwo.py" -input /map_reduce1/Result1.csv -output /map_reduce2
19/11/07 15:26:05 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.
packageJobJar: [/home/vaibhavoberoi7/MapperPhaseTwo.py, /home/vaibhavoberoi7/ReducerPhaseTwo.py] [/usr/lib/hadoop-mapreduce/hadoop-streaming-2.9.2.jar] /tmp/streamjob2954651502960744193.jar tmpDir=null
19/11/07 15:26:06 INFO client.RMProxy: Connecting to ResourceManager at cluster-2265-m/10.128.0.7:8032
```

Step 7 : We combine the 4 result files from Step 6, upload them to HDFS and run Phase Three.

Step 8 : Printing final results after combining the result files from Step 7.

```

vaibhavoeroi7@cluster-2265-m:~$ sed -e 's/\s/,/g' FinalResult.csv > FinalResult2.csv
vaibhavoeroi7@cluster-2265-m:~$ cat FinalResult2.csv
thelist,179736,0.00105584232755
gt,4883,0.00119808306709
gt,95592,0.00379106992418
gt,179736,0.00237564523698
gt,51816,0.00178997613365
gt,49153,0.00584320727155
gt,39677,0.00128929273084
gt,63051,0.000497347480106
gt,87234,0.000960717335611
Based,39677,0.00073673870334
Based,95592,0.00379106992418
consider,6068,0.0021645021645

```

Step 9: Now that we have got the results, we insert them into a table 'tfidf_results1', in order to find the terms and their respective tfidf scores for the top 10 users.

```

Time Taken: 0.824 seconds
hive> create table tfidf_results1(Term String, Id int, TFIDF float)
> row format delimited
> FIELDS TERMINATED BY ','
> location '/tf_idf'
> ;

```

OK
OK

```

1      0      4883      0.24937917
2      xargs  4883      0.04117109
3      checkbox 4883      0.03293687
4      idle   4883      0.024702653
5      branch 4883      0.024702653
6      blob   4883      0.024702653
7      tracking 4883      0.024702653
8      rm      4883      0.024702653
9      network 4883      0.024702653
10     composition 4883      0.024702653
1      0      6068      0.633416
2      assembly 6068      0.03304406
3      2px     6068      0.027036048
4      systemdatasqlite 6068      0.021028038
5      cultureneutral 6068      0.018024033
6      version10610 6068      0.018024033
7      incorrect 6068      0.018024033
8      publickeytokendb937bc2d44ff139 6068      0.018024033
9      modelstateisvalid 6068      0.018024033
10     sql      6068      0.017022697
1      practices 9951      0.03422053

```